

OffBoard product strategy, roadmap, and PRD

Executive summary

For this report, I use **OffBoard** as the future product name, while the current public repository and in-app branding are still **OffRecord**. The repo already contains a meaningful iOS-first foundation: a native SwiftUI/Xcode app, widget target, unit/UI tests, Core Data storage, optional CloudKit entitlements, biometric access, export services, and a local “Friday” assistant layer that performs sentiment/entity analysis and generates insights from journal data. The shipped roadmap also shows a clear ambition toward semantic search, Apple Foundation Models, and a more capable private assistant.

The market opportunity is real, but the category is split by trade-offs rather than by one clear winner. Day One has the strongest classic-journaling and encryption reputation, but its AI features can involve temporarily decrypting selected content and sending it to an AI service unless an on-device model is used. Rosebud and Reflection lean harder into AI guidance, weekly synthesis, and conversational support, but their public privacy models are less local-first than the posture you want for OffBoard. Daylio wins on speed, mood tracking, and local privacy, but not AI reflection. Apple Journal and Google Journal are powerful free defaults inside their ecosystems, with Apple emphasizing on-device suggestions and Google explicitly stating that Journal is not a substitute for therapy or professional mental health care.

The strongest strategic position for OffBoard is therefore **not** “another journaling app” and **not** “AI therapist.” It should be a **privacy-first, therapist-adjacent reflection companion**: voice-first capture, local-first storage and inference, careful weekly and monthly reflection, and explicit tools for preparing to talk with a real therapist, coach, or trusted person. That framing is also safer. Research on LLM-assisted journaling shows that AI can broaden reflection and help users see experiences from multiple perspectives, but it also warns about authorship loss, over-reliance, and undue trust. Separate mental-health-chatbot research is more cautionary still, especially for at-risk users, and major public guidance in health AI emphasizes careful limits, transparency, and human oversight.

The product recommendation is to launch **iOS-first and local-first**, because the repo already fits that path and Apple now provides an on-device foundation-model route that is tightly integrated with Swift and designed for private, offline inference on supported devices. Optional cloud AI should come later, behind explicit consent, narrow prompt scoping, and enterprise-grade retention controls such as OpenAI’s optional Zero Data Retention path or Anthropic’s ZDR/HIPAA-ready offerings where appropriate.

The single most important repo-audit finding is a **privacy-copy mismatch**. The README markets voice journaling as fully offline and says nothing is ever transmitted off-device, but the shipped `SpeechTranscriber` implementation says transcription is prioritized on-device and may fall back to Apple’s servers when online for better punctuation. That is fixable, but it means OffBoard’s privacy promise must be made more precise before scale.

Market landscape and positioning

Assumptions. Target geography, budget, and team size were not specified. I therefore assume a small product squad, consumer launch, and mobile-first development. Because the current codebase is a native Apple app with iPhone/iPad support, optional CloudKit sync, and an Apple-framework-heavy AI roadmap, the most efficient path is iOS-first rather than immediate parity across all platforms.

The market breaks into four practical segments: **ecosystem journals** that win through distribution, **classic journaling platforms** that win through memory-keeping and exportability, **mood and habit trackers** that win through low friction, and **AI-guided reflection apps** that win through follow-up questions, summaries, and emotional analysis. OffBoard should sit between classic privacy journaling and AI-guided reflection, while pushing harder than competitors on local-first architecture and therapist-prep workflows.

Segment	Primary user job	Representative products	Open gap for OffBoard	Sources
Ecosystem journals	"Give me a free, default journal tied to my phone and personal context."	Apple Journal, Google Journal	Stronger privacy controls, better exportability, and richer therapist-prep reflection than default apps offer today.	Official docs
Classic private journals	"Store my life safely, across devices, with durable exports."	Day One	More local-first AI and clearer processing-mode transparency.	Official docs
Mood and habit trackers	"Help me check in fast and show patterns."	Daylio	More narrative depth, voice capture, and guided review without losing low-friction check-ins.	Official docs
AI-guided reflection	"Talk to me, ask follow-ups, and summarize my patterns."	Rosebud, Reflection	Stronger local-first guarantees, clearer boundaries, and safer therapist-adjacent positioning.	Official docs

Competitor feature matrix

Product	Capture model	AI reflection depth	Reviews and insights	Privacy posture in public docs	Pricing and platforms	OffBoard takeaway	Sources
Day One	Text, media, audio, Apple Watch dictation, web/ browser access	Daily Chat, Go Deeper prompts, entry highlights, summaries, titles, image generation	Multi-entry summaries; strong export and memory-keeping posture	Core journals are E2EE; AI use requires consent and can temporarily bypass E2EE unless an on-device model is used	Basic free; Silver \$49.99/ year; Gold \$74.99/ year; iPhone, iPad, Android, Mac, Windows, web, Watch	Benchmark for premium journaling UX and export durability; beat it on default local AI and therapist-prep tools	Official docs
Rosebud	Text, voice journaling, call mode, custom journals	Ask Rosebud, dig deeper, entry reflection, long-term memory	Weekly report with themes, mood trends, people, and in-depth analysis after 1,500+ words/ week	Public privacy policy states collection of telemetry, event logs, model traces, performance metrics, and sharing with analytics/ service providers and anonymized research partners	Free + Bloom; site shows \$12.99/ month and discounted annual pricing	Users clearly value conversational support and weekly synthesis; OffBoard should offer similar value with materially stronger privacy guarantees	Official docs

Product	Capture model	AI reflection depth	Reviews and insights	Privacy posture in public docs	Pricing and platforms	OffBoard takeaway	Sources
Reflection	Text and voice across iOS, Android, Mac, web	AI Coach, AI writing support, AI search and insights	100+ expert-led guides; AI-enhanced search/insights; private entries can be excluded from AI	FAQ emphasizes AES-256 at rest and TLS in transit; changelog adds entry-level privacy controls	Premium at \$8/month or \$5.75/month billed annually; free trial	Reflection proves the market wants structured guidance plus “exclude this from AI” controls; OffBoard should copy that control, but local-first	Official docs
Daylio	Mood/activity-first with optional text, photos, and audio notes	Official positioning is not centered on generative coaching	Charts, goals, habit stats, “Year in Pixels,” exports	Strong local-first story: data stored locally, no data sent to Daylio servers, optional encrypted backup to user’s Google Drive	Free + premium; Android and iOS public offerings	Keep Daylio’s speed and stats discipline; narratively deepen it with better voice capture and review rituals	Official docs
Apple Journal	Text, photos, videos, locations, audio, state of mind, drawings/handwriting, share-sheet capture	Suggestions and reflection prompts based on on-device intelligence	Streaks, stats, insights, multiple journals, search, places view, export/print	Suggestions are on-device; App Store privacy label says identifiers and product-interaction analytics may be collected but not linked to the user	Free; iPhone, iPad, Mac within Apple ecosystem	Apple sets the baseline for polished default journaling; OffBoard must win on privacy transparency, voice-first capture, and deeper review	Official docs

Product	Capture model	AI reflection depth	Reviews and insights	Privacy posture in public docs	Pricing and platforms	OffBoard takeaway	Sources
Google Journal	Text, media, Gboard speech-to-text, optional cloud backup	Deep Dive, Revisit Topics, personalized reflections, Mood ID, AI insights	Goals, history, stats, insights	Google says content like inspirations and insights is processed on-device, but Gboard input may hit Google servers depending on settings; app is explicitly not a substitute for therapy or professional care	Free; Pixel 8+ app, AI features on Pixel 9+ except 9a; sign-in required for AI features	Best benchmark for explicit mental-health boundary setting; OffBoard should match or exceed that transparency	Official docs

The strategic implication is straightforward: **OffBoard should not chase maximum coaching breadth first**. It should launch with a narrower but more trustworthy promise: *“speak freely, keep everything private by default, get careful reflection and weekly synthesis, and bring selected insights to your therapist or trusted person if you choose.”* That promise is commercially differentiated because competitors usually optimize for either broad AI guidance or broad ecosystem convenience before they optimize for local control.

Users, value proposition, and success metrics

The best OffBoard users are people who already believe journaling could help them, but who hit one of four barriers: privacy anxiety, capture friction, lack of reflective structure, or inability to turn writing into patterns and conversations with real humans. That framing lines up with both HCI research on AI journaling and the feature patterns that have actually emerged across today’s leaders.

Persona	Core job to be done	What success looks like for them	Product metrics that matter most	Design implication
The Private Processor	“I need somewhere to say what I really think without being watched.”	Feels safe enough to write or speak honestly at least 3 times a week	Lock enablement, local-only mode retention, D30 retention, export confidence	Make local-only the default, show processing mode on every AI action, and make deletion/export obvious

Persona	Core job to be done	What success looks like for them	Product metrics that matter most	Design implication
The Voice-First Overthinker	"I can't keep up with my thoughts by typing."	Can capture a useful entry in under 90 seconds	Voice-to-save completion, transcription retry success, median capture time, draft recovery rate	Voice must be first-class, not a side feature; preserve raw audio until the user approves deletion
The Therapy Companion User	"Help me reflect between sessions and bring something useful to therapy."	Weekly reviews feel clarifying; therapy-prep summaries are actually shareable	Weekly review completion, summary export rate, "helped me prepare" rating, quote-redaction usage	Build explicit therapy-prep and selected-share workflows, not vague "AI coach" vibes
The Pattern Seeker	"Show me what keeps recurring in my life."	Sees meaningful trends and revisits them later	Search success, review revisit rate, knowledge-graph interaction, saved patterns/themes	Prioritize semantic search, recurring-theme detection, and reviews with evidence from entries

The recommended product promise is: **"OffBoard helps you capture your inner life, reflect with care, and notice patterns over time, without pretending to be your therapist and without treating your journal as training data."** That promise fits both the market gap and the current repo's natural strengths.

I recommend a north-star metric designed around **meaningful reflection**, not just opens or raw entry count. Journaling products fail when they optimize for vanity engagement or AI token volume; the better objective is repeated, user-owned reflection. DiaryMate is especially instructive here because it shows that AI can help users approach experiences from new perspectives, but can also erode authorship if the product over-optimizes machine contribution.

Metric layer	Metric	Definition	Why it matters
North star	Weekly Reflective Active Users	Unique users who complete at least 2 capture actions and 1 reflection action in a rolling 7-day window	Measures habit + reflection, not passive opens
Activation	First-session activation	Onboarding complete + privacy mode chosen + first journal saved within 24 hours	Ensures first-run UX is not too heavy
Habit	D7 / D30 retention	Users who return and save an entry in 7 / 30 days	Core PMF signal for journaling

Metric layer	Metric	Definition	Why it matters
Depth	Weekly review completion rate	Percent of eligible users who open and finish a weekly review	Confirms review value
Ownership guardrail	AI-dependence ratio	Share of final saved text generated by AI for users who use AI writing support	Should be monitored, not maximized
Privacy	Local-only retention	Share of actives who remain fully local and still retain	Tests whether local-first can sustain value
Safety	High-risk help engagement	Percent of high-risk interventions that lead to resource view, trusted-contact action, or self-declared safety-plan use	Measures whether safety UI is actionable
Business	Free-to-paid conversion	Share of monthly active users converting to paid plans	Tests monetization without harming trust

Roadmap and prioritized backlog

Planning assumptions. The timeline below assumes a small squad of roughly 4–6 people across product, design, mobile engineering, and backend/privacy support. Effort sizes use a simple t-shirt scale: **S** = around 1–2 engineer-weeks, **M** = 2–4 engineer-weeks, **L** = 1–2 engineer-months, **XL** = multi-quarter or cross-platform work.

A strong roadmap for OffBoard should deliberately sequence **trust before sophistication**. The MVP should not try to out-coach Rosebud or Reflection. It should instead make the privacy contract legible, make capture frictionless, make reviews genuinely useful, and make safety boundaries unmistakable. Only then should the product add semantic retrieval, encrypted sync, platform expansion, and premium cloud AI. That sequencing is also consistent with the repo’s present Apple-native shape and the current state of Apple/Android on-device model tooling.

Phase	Timeframe	Milestones	Effort	Dependencies	Key risks
MVP	0-3 months	Rebrand OffRecord → OffBoard; precise privacy copy; local-only onboarding; text/voice/photo capture; biometric vault; mood check-in; on-device AI reflection prompts; weekly review; export to PDF/Markdown/JSON; crisis-resource interruption UI; TestFlight beta	Mostly S-L	Privacy copy rewrite, transcription stabilization, review summarization, design system polish	Privacy mismatch, transcription failure, over-promising AI
V1	3-6 months	Semantic search; “private entry excluded from AI”; therapist-prep summary export; encrypted sync beta; reminders/goals refresh; payments and pricing tests; import from Apple/Day One/CSV where feasible	Mostly M-XL	Embeddings pipeline, sync key management, subscription plumbing	E2EE sync complexity, support burden, paywall timing
V2	6-12 months	Android pilot; knowledge graph and recurring-theme explorer; monthly/quarterly reviews; trusted contacts + safety plan; multilingual support; cloud AI add-on beta with strict consent and retention controls	Mostly L-XL	Android local inference, AI gateway, policy/eval stack, localization	Device fragmentation, privacy skepticism, safety false positives

Phase	Timeframe	Milestones	Effort	Dependencies	Key risks
V3	12+ months	Apple Foundation Models path on supported devices; Android LiteRT-LM path; cross-platform memory graph; watch/mac extensions; advanced evidence-cited review mode; selective care-team sharing	Mostly XL	Strong PMF, mature evaluation framework, sufficient budget	Anthropomorphism, regulatory scrutiny, feature sprawl

Prioritized backlog

Theme	Feature	Priority	Effort	Dependency	Why now
Capture	Text journaling with drafts/autosave	P0	S	Local store	Mandatory baseline
Capture	Voice capture with resilient save-before-transcribe flow	P0	M	Audio pipeline	Voice-first wedge
Capture	Photo attachments and inline media notes	P0	S	Local files	Already expected in category
Capture	Manual mood check-in plus optional AI mood estimate	P0	S	Entry schema	Enables reviews without forcing tagging
Capture	Search, filters, favorites, and date navigation	P0	M	Indexing	Core retrieval
Reflection	On-device follow-up prompts on selected text only	P0	M	Local inference	Core AI value without large trust cost
Reflection	Weekly review with themes, emotions, wins, friction, next questions	P0	M	Period summarizer	Key habit loop
Reflection	Therapist-prep summary export with quote redaction	P0	M	Export layer	Core therapist-adjacent differentiator

Theme	Feature	Priority	Effort	Dependency	Why now
Reflection	Semantic search over entries	P1	L	Embeddings	High-value retrieval
Reflection	Recurring-theme and people/topic knowledge graph	P1	L	Entity extraction	Pattern-recognition moat
Reflection	Monthly and quarterly life reviews	P2	M	Review engine	Long-term retention driver
Reflection	Evidence-cited reflection mode	P3	XL	RAG/eval stack	Trust and explainability moat
Privacy	Local-only mode with no account required	P0	S	None	Brand promise
Privacy	Processing-mode badge on every AI action	P0	S	UI framework	Trust legibility
Privacy	Export/delete everything from device	P0	S	Export layer	Ownership signal
Privacy	"Exclude this entry from AI/reviews/search" control	P1	M	Entry privacy tiers	Strong trust feature
Privacy	End-to-end encrypted sync	P1	XL	Key management	Cross-device expansion
Privacy	Minimal first-party analytics with no content collection	P0	M	Telemetry policy	Product learning without betrayal
Safety	Non-therapist disclaimers at onboarding and AI moments	P0	S	Legal copy	Regulatory and trust hygiene
Safety	Crisis keyword and semantic interruption router	P0	M	Safety rules	Must-have guardrail
Safety	Localized crisis resources and "contact trusted person" actions	P0	M	Resource dataset	Actionable escalation

Theme	Feature	Priority	Effort	Dependency	Why now
Safety	User-authored safety plan and trusted contacts	P1	M	Settings, share permissions	Better than generic resource screen
Safety	Red-team eval set and policy review workflow	P1	M	QA/ops	Keep model behavior bounded
Habit	Smart reminders and journaling windows	P1	S	Notification engine	Retention driver
Habit	Streaks, goals, and reflection streaks	P1	S	Goal model	Category baseline
Habit	Guided journal modes	P2	M	Prompt library	Broader use cases
Interop	Importers from Day One, Apple Journal, CSV/Markdown	P1	L	Parsing pipeline	Reduce switching cost
Interop	Share bundle for therapist, coach, or trusted partner	P1	M	Export layer	Therapist-adjacent moat
Platform	CI/CD for build, test, and release	P0	M	GitHub Actions/Fastlane	Repo maturity gap
Platform	Feature flags and remote copy controls	P1	M	Config service	Safer experimentation
Platform	Android app	P2	XL	PMF + Android inference	Important, but not first
Platform	macOS / watch extensions	P3	XL	PMF + Apple roadmap	Nice adjacency, not MVP
Monetization	Subscription/paywall framework	P1	M	Billing	Needed after value is clear
Monetization	Opt-in cloud AI add-on	P2	L	AI gateway	Premium tier lever

Product requirements document

Problem statement. People increasingly want journals that do more than store text: they want help naming emotions, noticing themes, generating weekly perspective, and preparing for real conversations with therapists or trusted people. But the current market forces uncomfortable trade-offs. The most private solutions are often weaker on active reflection, while the strongest AI-guided products generally require more cloud processing, more telemetry, or less legible data boundaries. The current OffRecord repo already offers a local-analysis foundation, but it does not yet provide the trust architecture, safety system, or therapist-adjacent workflows needed to make that positioning durable.

Product goal. Build a journaling app that helps users **capture, reflect, review, and share selected insights** while keeping privacy first and avoiding therapist impersonation. The product should help a user understand themselves better; it should not diagnose, treat, or claim clinical authority. That boundary is both strategically smart and aligned with public-health and regulatory guidance for low-risk wellness tools.

Goals. Make capture dramatically easier than typing. Make reflection more structured than a blank page. Preserve user authorship. Give users useful weekly and monthly reviews. Create explicit therapy-prep/share tools. Let users stay fully local by default.

Non-goals. OffBoard should not market itself as a therapist, mental-health professional, diagnostic system, suicide-prevention service, or emergency-response agent. It should not optimize for emotionally sticky AI companionship. It should not require cloud AI for core value.

Core user stories.

A user can create an entry by speaking, typing, or attaching media with no account. A user can request deeper reflection on selected text and clearly see whether processing is local or cloud-based. A user can complete a weekly review that summarizes recurring themes and cites source entries. A user can mark an entry as excluded from AI. A user who expresses high-risk language sees a safer interruption and direct paths to human help. A user can export their journal or a therapy-prep packet without handing over their full archive.

Acceptance criteria

Flow	Acceptance criteria
Onboarding	User can choose Local-only immediately; account creation is optional and deferred. App states clearly that it is not a therapist and cannot provide emergency care.
Capture	Voice note is saved locally before transcription; if transcription fails, the audio is preserved and the user can retry or type manually. Text and photo entries work offline.
AI reflection	AI only processes the current entry or user-selected excerpts by default; response is labeled Local or Cloud ; user can rate helpfulness or dismiss without saving.
Weekly review	Review covers a user-selected time range, summarizes themes/emotions, and links back to contributing entries or quotes. User can hide sensitive entries from the review and regenerate.
Safety escalation	When high-risk language is detected, the normal reflective AI flow is interrupted. The app shows supportive copy, crisis resources, trusted-contact options, and does not present itself as therapeutic advice.
Export and share	User can export full journal data in structured formats and create a separate, redacted summary bundle for therapy, coaching, or personal archival use.

The UX should deliberately preserve **user initiative**. DiaryMate found that people benefit from LLM-generated perspective, but also want control over tone, emotional valence, and the boundary between “my

writing” and “AI writing.” OffBoard should therefore treat AI as a **reflection scaffold**, not an autopilot co-author.

Key UX flows

flowchart TD

```
A[Open OffBoard] --> B[Choose Local-only or Optional Cloud AI Later]
B --> C[Set lock, reminder, and privacy preferences]
C --> D[Create entry]
D --> E{Capture mode}
E -->|Voice| F[Record and save audio locally]
E -->|Text| G[Write directly]
E -->|Photo| H[Attach media]

F --> I[Transcribe and save entry]
G --> I
H --> I

I --> J[Optional mood check-in]
J --> K[On-device safety triage]

K -->|Low or medium risk| L[Offer AI reflection on selected text]
L --> M[Prompts, reframes, and next questions]
M --> N[Save entry and reflection]

K -->|High-risk language| O[Interrupt reflective AI]
O --> P[Show human-help options and safety plan]

N --> Q[Weekly review queue]
Q --> R[Review themes, patterns, and quotes]
R --> S[Export full archive or shareable summary]
```

flowchart TD

```
A[Entry saved] --> B[Safety detector runs locally]
B --> C{Risk level}

C -->|None| D[No interruption]
C -->|Mild distress| E[Gentle grounding suggestions]
C -->|Moderate risk| F[Supportive note + resources + trusted contact option]
C -->|High risk| G[Stop normal AI coaching]

G --> H[Display crisis resources by locale]
H --> I[Offer call/text hotline]
H --> J[Offer message trusted contact]
H --> K[Open personal safety plan]
```

```

I --> L[User remains in control]
J --> L
K --> L

```

Data model

The current repo's persistence model only contains `DiaryEntry`, `PhotoAttachment`, and `AIState`. OffBoard needs a broader model to support privacy tiers, reviews, consent, safety, and selective sharing.

Entity	Key fields	Purpose
UserProfile	id, locale, timezone, age_gate_status, privacy_mode, reminder_prefs, biometrics_enabled	App-level preferences and safety boundaries
Journal	id, title, icon, color, created_at, sync_enabled, ai_enabled	Multiple journals/spaces
Entry	id, journal_id, created_at, updated_at, source_type, raw_text, cleaned_text, manual_mood, ai_mood, starred, privacy_tier, risk_level	Core journal record
EntryAttachment	id, entry_id, type, local_uri, mime_type, duration, created_at	Audio, photo, file attachments
ReflectionSession	id, entry_id, mode, provider, processing_mode, prompt_scope, output_text, saved_to_entry, helpfulness_rating	AI reflection interactions
ReviewReport	id, journal_id, period_start, period_end, summary, key_themes, emotional_trend, wins, frictions, quote_refs, generated_at	Weekly/monthly reviews
KnowledgeNode	id, label, type, first_seen, last_seen, mention_count, sentiment_weight	People/topics/places for patterning
KnowledgeEdge	id, from_node, to_node, relation_type, weight, last_seen	Lightweight graph structure
Goal	id, journal_id, type, target_value, cadence, status	Entries-per-week or reflection habits
ConsentEvent	id, user_id, consent_type, scope, granted, timestamp, policy_version	Auditable consent record
SafetyEvent	id, entry_id, detector_version, labels, severity, intervention_type, resource_locale, action_taken	Safety logging without storing excess content
ShareBundle	id, review_id, bundle_type, redaction_level, recipient_type, expires_at	Therapy-prep or trusted-contact sharing

Entity	Key fields	Purpose
SyncEnvelope	id, entity_type, entity_id, ciphertext_blob, wrapped_key_id, device_id, updated_at	E2EE sync unit
SubscriptionState	id, plan, started_at, renewal_at, ai_credits_remaining	Monetization layer

```

erDiagram
    UserProfile ||--o{ Journal : owns
    Journal ||--o{ Entry : contains
    Entry ||--o{ EntryAttachment : has
    Entry ||--o{ ReflectionSession : triggers
    Journal ||--o{ ReviewReport : generates
    Journal ||--o{ Goal : tracks
    Entry ||--o{ SafetyEvent : may_create
    ReviewReport ||--o{ ShareBundle : exports
    Journal ||--o{ SyncEnvelope : syncs
    UserProfile ||--o{ ConsentEvent : records
    Journal ||--o{ KnowledgeNode : derives
    KnowledgeNode ||--o{ KnowledgeEdge : links

```

API and infrastructure needs

The current repo is fundamentally a local Apple app with optional CloudKit intent, not a service-oriented multiplatform system. For OffBoard, the infrastructure should remain intentionally thin at first.

Capability	MVP recommendation	Later recommendation
Local persistence	Keep local-first encrypted store and file vault	Abstract into a shared domain layer for iOS/Android
On-device AI	Local rules + small summarization/prompting stack	Multiple local runtimes by platform
Sync	No mandatory account; optional encrypted sync beta	Cross-platform E2EE sync service with ciphertext-only storage
Cloud AI gateway	Not required for MVP	Narrowly scoped gateway that sends only selected excerpts, strips metadata, and honors org-level retention controls
Search/index	On-device lexical search first	Add embeddings and local vector index
Notifications	On-device reminders only	Optional cross-device reminder state
Analytics	Minimal first-party telemetry, zero content	Cohort/event analytics with strong privacy defaults

Capability	MVP recommendation	Later recommendation
Safety services	Local risk rules and static resource bundles	Server-updated policies/resources only if users opt in

Privacy and security design. OffBoard should use a **local-first default**, a **no-account-required core experience**, and **processing-mode transparency** on every AI feature. Optional sync should be end-to-end encrypted with wrapped keys and ciphertext-only server storage. Optional cloud AI should be strictly **per-request**, never full-journal by default, and routed through a gateway that sends only user-selected text. For vendor processing, the strongest posture is to use enterprise controls where content is not used for training by default and retention is minimized or formally zero-retention where supported. Reflection's recent "private entries excluded from AI" feature is an especially good pattern to adopt early.

Compliance and regulatory considerations. OffBoard should be designed and marketed as a **low-risk wellness product**, not a diagnosis or treatment tool. In the United States, HIPAA may not apply to a direct-to-consumer journaling app unless it is acting for a covered entity or business associate, but the FTC's Health Breach Notification Rule can apply to vendors of personal health records and related entities outside HIPAA. In the EU, mental-wellbeing and health-adjacent journaling data is likely to trigger special-category-data considerations under GDPR Article 9 and should be treated accordingly, including explicit consent, strict purpose limitation, and a DPIA before scaling cloud AI features.

Safety guardrails. The app should use layered language that says what it *is* and what it *is not*: a personal reflection tool, not a therapist; a helpful summarizer, not a diagnostic authority; a journaling companion, not a crisis responder. Google's Journal product makes this distinction explicitly, and research on mental-health chatbots strongly supports that caution. OffBoard should use local risk detection first, interrupt normal reflective AI when necessary, and guide users toward humans, crisis lines, or trusted contacts without claiming to manage emergencies itself.

Technical audit of the current repo

The repo is significantly more than a landing page or prototype. It already contains a functioning Apple-native journal product with multiple app targets, export logic, analytics views, app-store assets, documentation, and a visible future roadmap. But it is still pre-scale: the current public repo shows no published releases, only one obvious GitHub Pages workflow, and multiple signs of branding and policy debt that should be cleaned up before broader launch.

Audit finding	Current state	Product implication	Evidence
Product scope	Native SwiftUI/Xcode project with app, widget, tests, marketing, App Store assets, website, docs	Strong iOS-first base; not yet a true multiplatform product	Repo root
Core storage	<code>DiaryEntry</code> , <code>PhotoAttachment</code> , <code>AIState</code> in Core Data	Needs schema expansion for consent, safety, reviews, sharing, subscriptions	Data model

Audit finding	Current state	Product implication	Evidence
AI today	Local AI engine + Friday assistant + Friday chat + insights views	Good local-analysis base, but currently closer to structured NLP plus template logic than a robust assistant platform	Code and roadmap
Security/export	AES-256-GCM encrypted backups, JSON/Markdown/CSV/PDF export, Face ID, iCloud entitlements	Excellent ownership/privacy starting point	Code and plist/entitlements
Privacy-copy mismatch	README says everything is offline and never leaves the device; speech transcriber may use Apple servers when online	Must fix copy or remove fallback to prevent user-trust damage	README + transcription code
CI/CD maturity	One GitHub Pages workflow deploying <code>website/public</code> ; no public evidence of app build/test/release pipelines	Add automated testing, linting, build, signing, and release process before scale	Workflow + root listing
Branding consistency	Placeholder <code>offrecord.example.com</code> links in README; export text includes "DAILYVOX DIARY EXPORT"	Clean branding is required before monetization or press	README + backup service
Roadmap ambition	Future plans include semantic search, Apple Foundation Models, LoRA fine-tuning, Personal Voice, macOS/Watch	Vision is strong, but sequencing needs to be tightened around trust/safety first	Roadmap

Missing components to add before OffBoard launch. The repo does not yet show a full safety escalation system, consent ledger, privacy dashboard, entry-level AI exclusion control, encrypted multiplatform sync architecture, cloud-AI provider abstraction, feature flags, mature observability, or publication-ready release automation. Some of these are absent by omission; others are explicitly listed as future roadmap items rather than shipped capabilities.

Recommended repo structure. I would move to a product-oriented monorepo once OffBoard goes beyond the current Apple-only shape:

```
offboard/
  apps/
    ios/
    android/
    web/
  packages/
    domain/
```

```

design-system/
crypto/
local-ai/
search/
safety/
exports/
services/
  sync-api/
  ai-gateway/
  notifications/
  billing/
infra/
  terraform/
  policies/
  ci/
docs/
  prd/
  adrs/
  privacy/
  safety/
ml/
  evals/
  prompts/
  on-device-models/

```

Tech stack recommendation. For MVP, the highest-leverage path is to **keep iOS/iPad native** and aggressively modularize the existing SwiftUI codebase, because Apple’s on-device model stack is now directly integrated with Swift and the repo already assumes that world. If Android parity becomes a must before PMF, a cross-platform UI framework can help, but both Flutter and React Native would still need deep native bridges for biometrics, encrypted local file handling, widgets, selective AI processing, and Apple/Android on-device inference frameworks. In other words, a cross-platform UI does **not** remove the hardest engineering problems in this category.

On-device and cloud AI options

Option	Best fit for OffBoard	Integration point	Privacy and product trade-off	Sources
Apple Foundation Models	iOS-first journaling prompts, structured weekly reviews, light conversational reflection on supported Apple devices	Native Swift modules in iOS app	Best privacy story: on-device, offline, no added app size, tightly aligned with Swift; device-gated to Apple Intelligence hardware	Apple docs and WWDC/ materials

Option	Best fit for OffBoard	Integration point	Privacy and product trade-off	Sources
Core ML custom models	On-device classification, embeddings, mood/risk heuristics, lightweight summarizers	Native local inference layer	Great for privacy and latency, but model quality and maintenance are your burden	Apple docs
Android LiteRT-LM / MediaPipe path	Android V2 for on-device summaries, retrieval prompts, and reflection assists	Android inference module	Strong on-device option, but optimized for high-end devices and models are too large to bundle directly in an APK	Google AI Edge docs
llama.cpp + GGUF models	Cross-platform custom local model experimentation	Shared local runtime/native modules	Flexible and open, but heavier battery/performance/app-size work than platform-native frameworks	Official repo
Meta Llama 2-family models	Candidate model family for self-hosted/local experimentation where licensing is acceptable	Via Core ML conversion or llama.cpp	Commercial use is possible under Meta's community license, but review the license and acceptable use terms carefully	License source
Stanford Alpaca	Do not use in production	N/A	Research-only and non-commercial; unsuitable for a consumer product	Stanford sources
OpenAI API	Opt-in premium cloud reflection, long-context review generation, multilingual polish	Cloud AI gateway only	API data is not used for training by default, but abuse logs are retained up to 30 days unless Zero Data Retention or modified controls are approved	Official docs
Anthropic Claude API	Opt-in premium cloud reflection and summary generation with stronger enterprise/privacy controls	Cloud AI gateway only	Offers ZDR for eligible API usage and HIPAA-ready API access with a signed BAA; still requires explicit scope and governance	Official docs

Recommendation. For OffBoard, use a **tiered inference strategy**: local rules + small on-device models for MVP, Apple Foundation Models for supported iOS devices as soon as they fit the product, Android LiteRT-LM for later Android parity, and premium cloud AI only as an opt-in layer for selected workloads that genuinely

need longer context or stronger generation quality. That preserves your differentiation instead of eroding it.

Monetization, go-to-market, and first-year KPIs

The current repo markets OffRecord as **free forever**, with no subscriptions, no in-app purchases, and no ads. That promise is powerful, but it also limits product investment unless OffBoard creates a monetization layer that does **not** violate the privacy-first brand. Competitor anchors suggest room for premium pricing if the value is clear: Day One Gold is \$74.99/year, Reflection promotes roughly \$8/month or \$5.75/month billed annually, and Rosebud Bloom publicly shows \$12.99/month with discounted annual pricing.

Monetization options and pricing experiments

Experiment	Offer	Hypothesis
Free Local	Unlimited local journaling, local AI reflection, weekly review, exports, no account required	A generous free local tier is necessary because Apple and Google provide free defaults
Privacy Plus	~\$49–59/year for encrypted sync, multiple journals, advanced exports/imports, private-entry vault, reminder packs	Users will pay for ownership and cross-device continuity without needing cloud AI
Therapy Companion	Add therapist-prep summaries, recurring-theme packets, and trusted-share bundles to paid tier	The therapist-adjacent value prop is monetizable and differentiated
Cloud AI Add-on	~\$6–8/month or credits for long-context summaries, multilingual support, premium review modes	Users will accept paid cloud AI if it is narrow, explicit, and optional
Annual-first trial	7–14 day free trial on annual plan vs no-trial freemium gate	Annual plans should outperform monthly if trust and privacy value are explicit
Power-user bundle	Bundle Privacy Plus + Cloud AI below Rosebud but near Day One Gold	Mid-market premium positioning may win users who reject both cheap/coarse and expensive/cloud-heavy options

The go-to-market should emphasize **trust language users can verify**. That means product pages that explain local-only mode, processing-mode labels, export/delete controls, and what exactly happens when cloud AI is enabled. The best early wedge is likely privacy-conscious iPhone users, voice-first diarists, therapists/coaches who want a between-sessions companion for clients, and users dissatisfied with either blank-page journals or chatty AI companions. Competitor evidence already shows demand for weekly summaries, AI follow-up, and therapist-adjacent use, but it also shows that privacy posture is still a messy, under-served part of the category.

The most effective GTM motions are likely:

- an **iOS-first launch** into journaling, productivity, ADHD, and privacy communities;
- landing-page messaging around **“AI journaling that stays on your device”**;
- content/SEO around **therapy prep, voice journaling, and private reflection**;
- selected pilots with therapists, coaches, and university counseling-adjacent wellness groups, while staying careful not to market the app as treatment.

First-year KPI plan

Period	Product KPI targets	Business KPI targets	Risk KPI targets
Months 1-3	60% first-session activation; D7 retention > 28%; weekly review completion > 25% of eligible users	No monetization required; focus on waitlist-to-activation > 35%	< 1 privacy trust incident per 1,000 actives; crash-free sessions > 99.5%
Months 4-6	D30 retention > 18%; Weekly Reflective Active Users / MAU > 32%; AI helpfulness > 4.2/5	Free-to-paid conversion 3-5%; annual plan mix > 50% of paid starts	High-risk resource engagement > 35% of flagged sessions
Months 7-12	D30 retention > 24%; WRAU / MAU > 38%; semantic search usage among actives > 20%	Free-to-paid conversion 6-8%; paid churn < 4.5% monthly	< 0.5% of support tickets tied to privacy confusion; zero severe safety policy violations

Prioritized sources

The most decision-worthy sources for this project are the ones below, because they directly anchor product scope, competitor reality, privacy controls, and safety/regulatory constraints.

- **Current OffRecord codebase and docs** for shipped scope, roadmap, data model, exports, permissions, and repo maturity.
- **Day One official docs** for E2EE, AI-processing trade-offs, and pricing tiers.
- **Rosebud official site/help/privacy** for pricing, weekly-report design, long-term-memory positioning, and telemetry/privacy posture.
- **Reflection official site/FAQ/changelog** for AI-coach positioning, 100+ guide library, encrypted storage claims, private-entry controls, and premium pricing.
- **Apple and Google official journal materials** for ecosystem-default feature baselines and explicit “not a therapist” boundary language.
- **Academic/HCI literature** on AI-supported journaling and mental-health-chatbot risks, especially DiaryMate, MindScape, and the 2025 JMIR mental-health-chatbot analysis.
- **Privacy and compliance sources** from HHS, FTC, FDA, and the European Commission/GDPR for legal framing.
- **Official AI platform docs** from Apple, Google, OpenAI, and Anthropic for on-device versus cloud-AI architecture choices and retention trade-offs.